AuriVault Brand + Product Blueprint

🔑 Brand Positioning

Essence: Secure brilliance — turning hidden organizational knowledge into radiant, actionable insight.
Core Promise: AuriVault protects knowledge like a vault while illuminating it like gold.

Tagline Options: - *"Unlock the Gold in Your Data."*
- *"Where Knowledge Glows, Securely."*
- *"Illuminate. Secure. Transform."*

🕐 Visual Identity

Colors: - **Deep Midnight Blue** → Trust, stability, security.
- **Gold Accents** → Illumination, premium quality.
- **White** → Clarity, simplicity, openness.

Logo Concept: A vault door morphing into a radiant sunburst — merging the imagery of protection and enlightenment.

Design Language: Minimal, clean dashboards with radiant highlights (golden glows on call-to-actions, graphs, or insights). Typography should balance modern sans-serif clarity with subtle serif elegance for premium positioning.

📖 Narrative

AuriVault isn't just another knowledge repository.
It is a **living intelligence layer** that: - Transforms static documents into **active, glowing insights**. - Protects organizational memory with the **strength of a vault**. - Illuminates knowledge gaps, enabling teams to **anticipate, not just react**. - Evolves with each tenant's unique context, ensuring **personalized brilliance**.

Tone of Voice: Authoritative yet radiant, combining trustworthiness with inspiration.

🗜️ Product Anchoring

Directly mapping brand essence to SaaS differentiators:

- **Vault (Security & Trust)**
- Multi-tenant isolation
- AES-256 encryption at rest, TLS in transit
- Role-based access control

- Audit logs & GDPR compliance

- **Aura (Illumination & Insight)**

- Tailored embeddings & fine-tuning per tenant
- Gap analytics: unanswerable queries, recurring clusters, missing docs
- Proactive recommendations for document updates
- Transparent AI with citations & confidence scoring

🧰Core Features

User Experience - Secure login & workspace isolation.
- Document uploads (PDF, DOCX, TXT, CSV, images with OCR).
- Chat assistant with contextual, cited responses.
- Dashboard with query history & usage metrics.

AI Features - Embeddings: Instructor, MiniLM, OpenAI embeddings (configurable).
- Retrieval-Augmented Generation (RAG) pipeline with citations.
- Fine-tuning: LoRA/PEFT adapters for per-tenant models.

Analytics & Insights - Top queries per week.
- Confidence scoring for responses.
- Knowledge gap identification.
- Role-based usage analytics.

Integrations - Slack, Teams, LMS chatbot connectors.
- REST API & webhooks.
- Exportable reports (CSV, PDF).

🏦System Architecture (Conceptual)

1. **User / Tenant** → Auth & Multi-Tenancy Layer
2. **File Upload Module** → Text extraction & OCR
3. **Document Processing** → Chunking & metadata tagging
4. **Embeddings Module** → Chunks → Vector Store (pgvector, Qdrant, Neon)
5. **RAG Query Pipeline** → Retrieve docs → LLM inference → Response with citations
6. **Chat UI / API Output**

Parallel Paths: - Analytics Engine (dashboards & insights)
- Fine-Tuning Module (per-tenant adapters)
- Security & Compliance (audit logs, encryption)

🔒 Security & Compliance

- Row-level tenant isolation.
- Encrypted storage (AES-256).

- TLS encryption in transit.
- Role-based access control.
- Audit logs for queries & responses.
- GDPR-aligned data handling.

🛠️ Tech Stack

Frontend: Next.js + Tailwind, tRPC for type-safe API.
Backend: Node.js/TypeScript (T3 stack), Prisma ORM, PostgreSQL/Neon.
Vector DB: pgvector, Qdrant, or Weaviate.
AI/ML: Instructor-XL, MiniLM, OpenAI embeddings; LLaMA-2/3, Mistral; LoRA adapters.
Deployment: Vercel (frontend), Supabase/Neon (DB/storage), Render/Fly.io (backend).
Monitoring: Latency, accuracy, error rates, feedback logging.

📈 Roadmap

Phase 1 – MVP (4–6 Weeks) - Multi-tenant login & workspaces.
- Document upload & processing.
- Embedding & retrieval with vector store.
- Basic chat assistant with citations.
- Minimal dashboard (query history).

Phase 2 – Analytics (4 Weeks) - Usage statistics & top queries.
- Low-confidence query detection.
- Gap analysis & recommendations.
- Exportable CSV/PDF reports.

Phase 3 – Personalization (6 Weeks) - Tenant-specific LoRA fine-tuning.
- Multiple models per tenant.
- Switch between default & tenant-trained model.

Phase 4 – Integrations (6 Weeks) - Slack/Teams/LMS connectors.
- REST API & webhooks.

Phase 5 – Scaling & Security (Ongoing) - Audit logs & compliance enhancements.
- Multi-region deployment.
- Billing & subscription plans.

🧱 Differentiators for Clients

- **Tailored AI:** Each tenant gets a personalized assistant.
- **Actionable Analytics:** Illuminate gaps, not just answers.
- **Transparency:** Every response cites sources.
- **Cost-Efficiency:** Open-source first, scale when needed.
- **Integration-Friendly:** APIs & connectors for workplace adoption.

📌 Next Steps

1. Finalize brand identity (logo, tagline, palette).
2. Secure domain + trademark.
3. Set up repo & skeleton (create-t3-app).
4. Build MVP (Phase 1).
5. Onboard pilot users.

🕐 AuriVault Promise

"AuriVault is where security meets illumination — a golden vault of living knowledge, empowering organizations to protect, unlock, and transform their most valuable asset: insight."